



PASS
SUMMIT 2016

Lockless in Seattle

Using In-Memory OLTP for
Transaction Processing

Mark Broadbent, Principal,
SQLCloud





Please silence
cell phones

A movie poster background featuring a man and a woman standing on a rooftop in Seattle. The man, on the left, is wearing a grey sweatshirt and dark pants, looking up at the sky with his arms crossed. The woman, on the right, is wearing a light-colored blazer and dark pants, looking up at the sky with her hand on a railing. The background shows a city skyline at sunset or sunrise, with a full moon and stars visible in the dark sky.

MARK BROADBENT

LOCKLESS

IN SEATTLE

Using In-Memory OLTP
for Transaction Processing

Principal Consultant

The logo for sqlcloud, featuring the word "sql" in blue and "cloud" in a lighter blue, with a stylized cloud icon above the "d".

SQLCLOUD.CO.UK


Explore Everything PASS Has to Offer



FREE ONLINE WEBINAR EVENTS



FREE 1-DAY LOCAL TRAINING EVENTS



VOLUNTEERING OPPORTUNITIES



**LOCAL USER GROUPS
AROUND THE WORLD**



**ONLINE SPECIAL INTEREST
USER GROUPS**



PASS COMMUNITY NEWSLETTER

PASS BLOG

WHITE PAPERS

SESSION RECORDINGS



FREE ONLINE RESOURCES



BUSINESS ANALYTICS TRAINING



BA INSIGHTS NEWSLETTER

PASS VIRTUAL CHAPTERS

www.sqlpass.org/vc

PASS APPLICATION
DEVELOPMENT
VIRTUAL CHAPTER 

PASS BIG DATA
VIRTUAL CHAPTER 

PASS BUSINESS
ANALYTICS
VIRTUAL CHAPTER 

PASS BUSINESS
INTELLIGENCE
VIRTUAL CHAPTER 

PASS CLOUD
VIRTUAL CHAPTER 

PASS DATA
ARCHITECTURE
VIRTUAL CHAPTER 

PASS DATABASE
ADMINISTRATION
VIRTUAL CHAPTER 

PASS DATA
SCIENCE
VIRTUAL CHAPTER 

PASS DBA
FUNDAMENTALS
VIRTUAL CHAPTER 

PASS EXCEL BI
VIRTUAL CHAPTER 

PASS GLOBAL
CHINESE
VIRTUAL CHAPTER 

PASS GLOBAL
FRANCOPHONE
VIRTUAL CHAPTER 

PASS GLOBAL
HEBREW
VIRTUAL CHAPTER 

PASS GLOBAL
ITALIAN
VIRTUAL CHAPTER 

PASS GLOBAL
PERSIAN
VIRTUAL CHAPTER 

PASS GLOBAL
PORTUGUESE
VIRTUAL CHAPTER 

PASS GLOBAL
RUSSIAN
VIRTUAL CHAPTER 

PASS GLOBAL
SPANISH
VIRTUAL CHAPTER 

PASS HEALTHCARE
VIRTUAL CHAPTER 

PASS HIGH AVAILABILITY &
DISASTER RECOVERY
VIRTUAL CHAPTER 

PASS HYBRID
VIRTUAL CHAPTER 

PASS INFRASTRUCTURE
VIRTUAL CHAPTER 

PASS IN MEMORY
VIRTUAL CHAPTER 

PASS PERFORMANCE
VIRTUAL CHAPTER 

PASS POWERSHELL
VIRTUAL CHAPTER 

PASS PROFESSIONAL
DEVELOPMENT
VIRTUAL CHAPTER 

PASS SATURDAY
NIGHT SQL
VIRTUAL CHAPTER 

PASS SECURITY
VIRTUAL CHAPTER 

PASS VIRTUALIZATION
VIRTUAL CHAPTER 

PASS WOMEN
IN TECHNOLOGY
VIRTUAL CHAPTER 

Session Evaluations

Your feedback is
important and valuable.

3

ways to access

Submit by 5pm
Friday November 6th to
WIN prizes



Go to passSummit.com



Download the GuideBook App
and search: **PASS Summit
2016**



Follow the QR code link displayed
on session signage throughout the
conference venue and in the
program guide

Contact...



mark.broadbent@sqlcambs.org.uk



@retracement



tenbulls.co.uk

Likes...



Guilty pleasures...



Badges...

**Microsoft
CERTIFIED**

Master: SQL Server




Community...



SQLA



Agenda



1

Why IMOLTP?



2

Architecture



3

Implementation



4

Isolation & TP
Control



4

Limitations

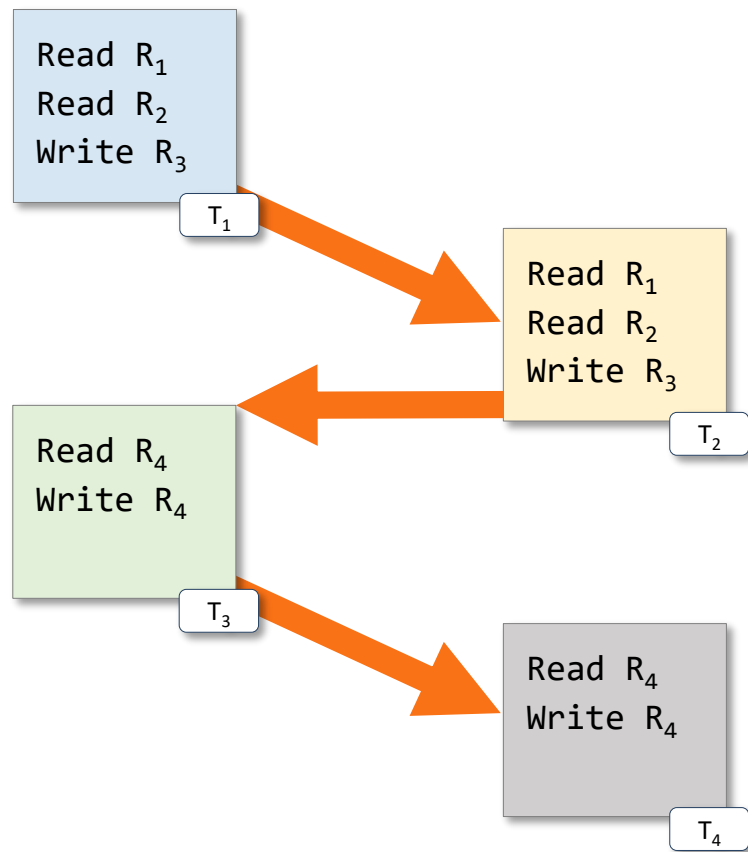


5

The best part of the
presentation... Gin O'Clock

Serial Processing

- Theoretically at least, serial execution time (should) = our slowest throughput speed.
(This is the 2nd law of Concurrency Control)
- Our aim then is to execute workloads in parallel right?!
- But when we do, our typical resource bottlenecks are:
 - Memory
 - CPU
 - Disk IO



Parallel Processing requires Transaction Interleaving

T₂ (write) waits for T₁ to
release Slock on R₁
Slock is released when?

Lock
wait

Pessimistic

T ₁	Read	R ₁
T ₂	Write	R ₁
T ₂	Read	R ₂
T ₁	Write	R ₃
T ₂	Write	R ₃
T ₃	Read	R ₄
T ₃	Write	R ₄
T ₄	Read	R ₄
T ₄	Write	R ₃

T₂ (write) waits for T₁ to
release Xlock on R₃
Xlock is released when?

Lock
wait

Disk based
Optimistic

T₂ (write) is not blocked
by T₁ (read)
Why no blocking?

T₂ (write) waits for T₁ (write)
to release Xlock on R₃
Why is there blocking?

Lock
wait

Parallel Processing requires Transaction Interleaving

T_2 (write) will not be
blocked by T_1 (read)
(same behaviour on-disk)

IMOLTP
(Optimistic)

T_2 (write) will not be
blocked by T_1 (write)
Why are writes not blocked?

T_1	Read	R_1
T_2	Write	R_1
T_2	Read	R_2
T_1	Write	R_3
T_2	Write	R_3
T_3	Read	R_4
T_3	Write	R_4
T_4	Read	R_4
T_4	Write	R_3

Concurrency Models in SQL 2014 and beyond

- Pessimistic Isolation

- Readers do not block readers
- **Writers block readers**
- **Readers block writers**
- **Writers block writers**

- (disk based) Optimistic Isolation


- Readers do not block readers
- Writers do not block readers
- Readers do not block writers
- **Writers block writers**

- (In-Memory) Optimistic Isolation

- Readers do not block readers
- Writers do not block readers
- Readers do not block writers
- Writers do not block writers

Governed by lock compatibility and implemented through lock hierarchy and escalation



Governed by write conflict detection



Sch-M, or object level X can kill concurrency.
Intent locks and escalation have overhead!

Hardware Trends

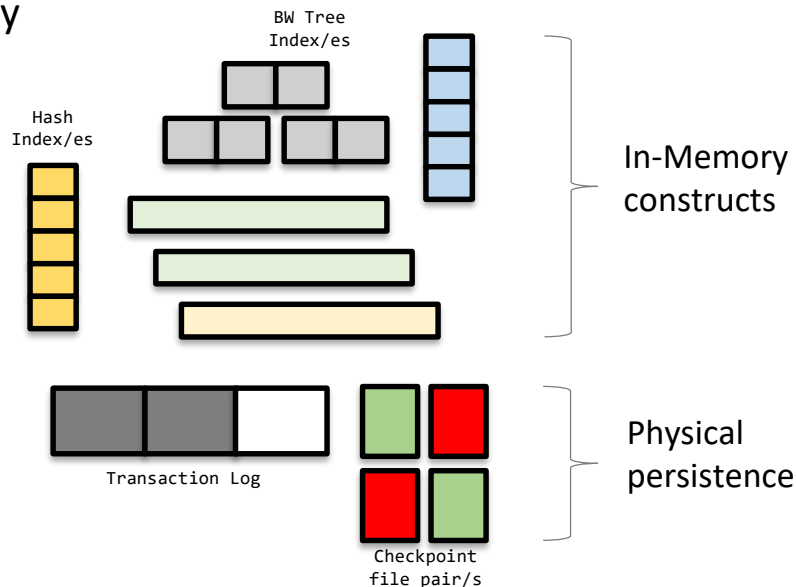
- Hardware trend is for scale not speed
- CPU Core count increases, clock speed static
- Memory sizes increase/ costs fall
- (As disk speeds also increase)
- Concurrency is clearly a software problem



Graphics from BRK3576, In-Memory – The Road Ahead by Kevin Farlee – Ignite Conference 2015

In-Memory OLTP to the rescue...


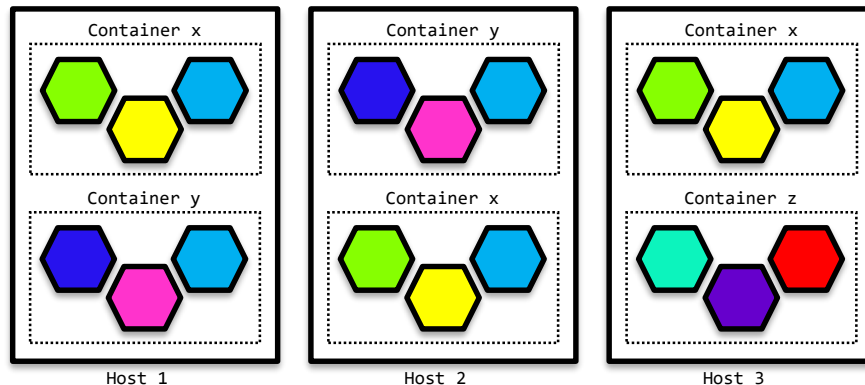
- In-Memory data structures (optimised for memory)
- Persistence through Transaction Log and checkpoint file pair/s
- Logging optimizations and improvements
- No TempDB overhead – all versioning In-Memory
- Lockless and latchless operation
- No fragmentation concerns
- Baked into product



A solution to a problem that doesn't exist?

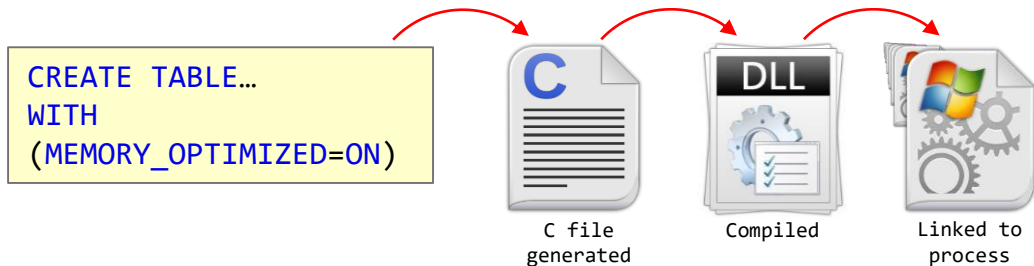
- Before Internet we were in db-centric world. We are now in a message centric world
 - A defined platform tends to define a solution, when in fact it should be the problem driving the solution...
 - Micro-services, stateless data queues, smaller data repositories, no locking
 - Application Caching Fabrics reduce latencies
- ***Newsflash*** so does SQL Server Caching/ Buffering

- Are monolithic OLTP environments about to become extinct?




In-Memory Data Structures (Tables and Indexes)

- Tables
 - On create - C file generated, Compiled to DLL and Linked to SQL process
 - Structure of the row is optimized for memory residency and access
 - Recreated, compiled, linked on database startup
- Indexes (rebuilt at startup)
 - Hash indexes for point lookups
 - Memory-optimized non-clustered index for range and ordered scans
 - Do not duplicate data, just pointers on rows
 - **Warning! Statistics not auto-updated.**



In-Memory Data Structures (Natively Compiled Stored Procedures)

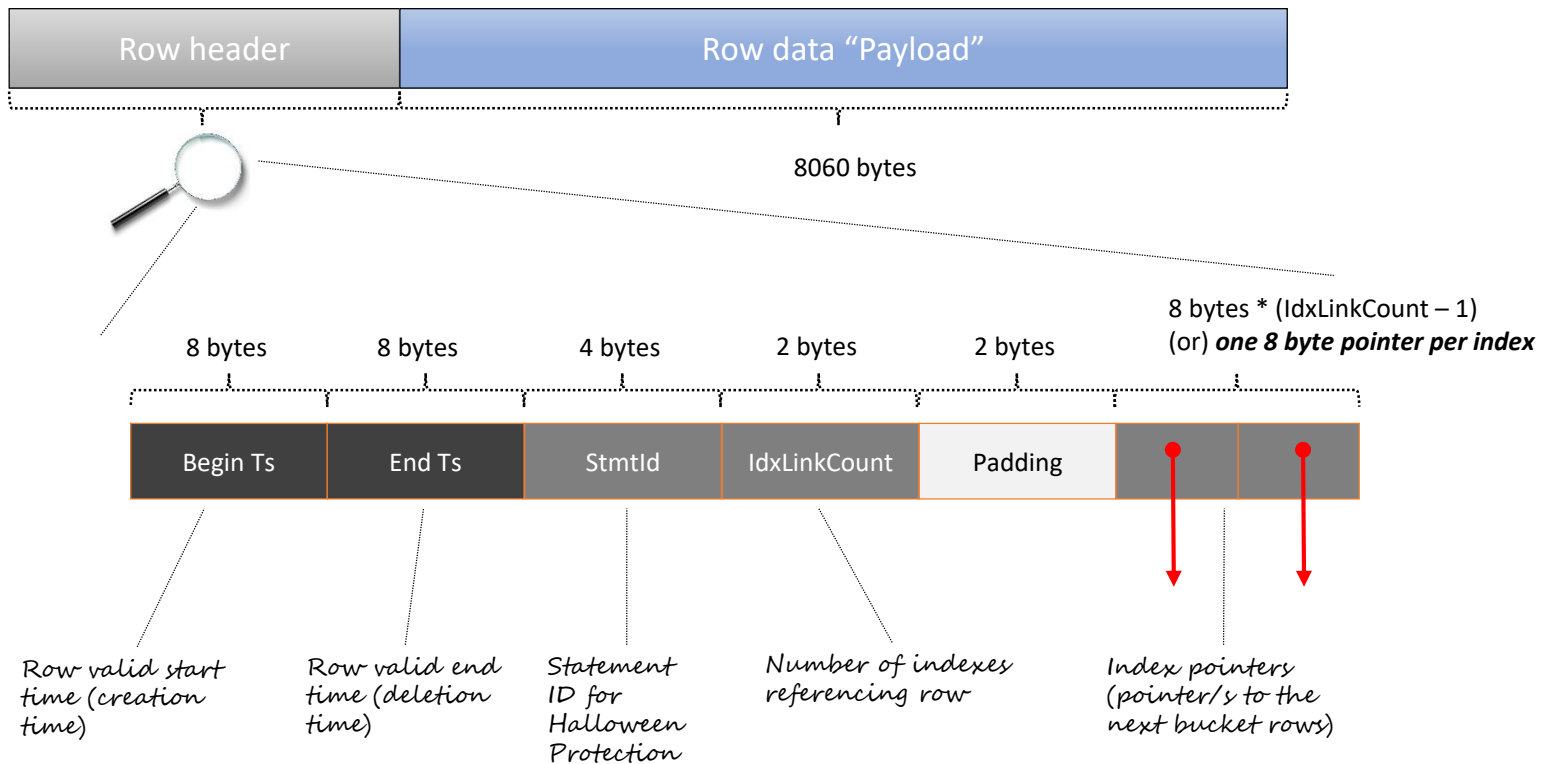
- On create - Optimized, C file generated, Compiled to DLL and Linked to SQL process
- **Warning, optimised once and Plan based on statistics!**
- Not part of plan cache (so not visible in `sys.dm_exec_cached_plans`)
- Recreated on database startup and compiled on first execution
- They can only access In-Memory Tables



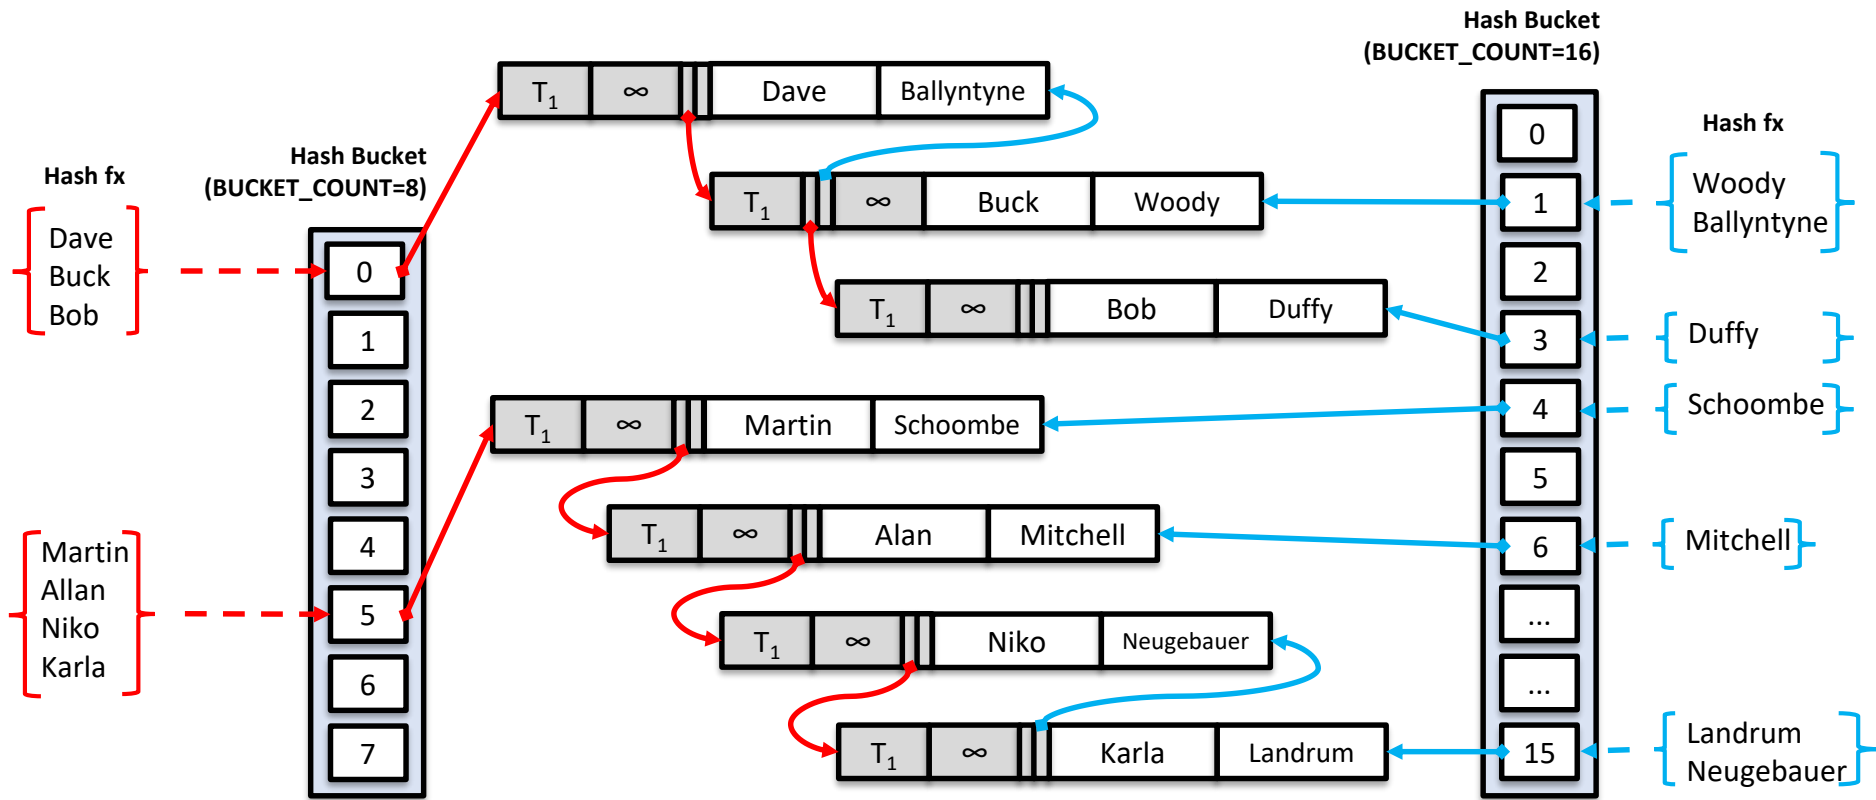
Demo

Creating in memory tables

MemOpt Table: Row Format



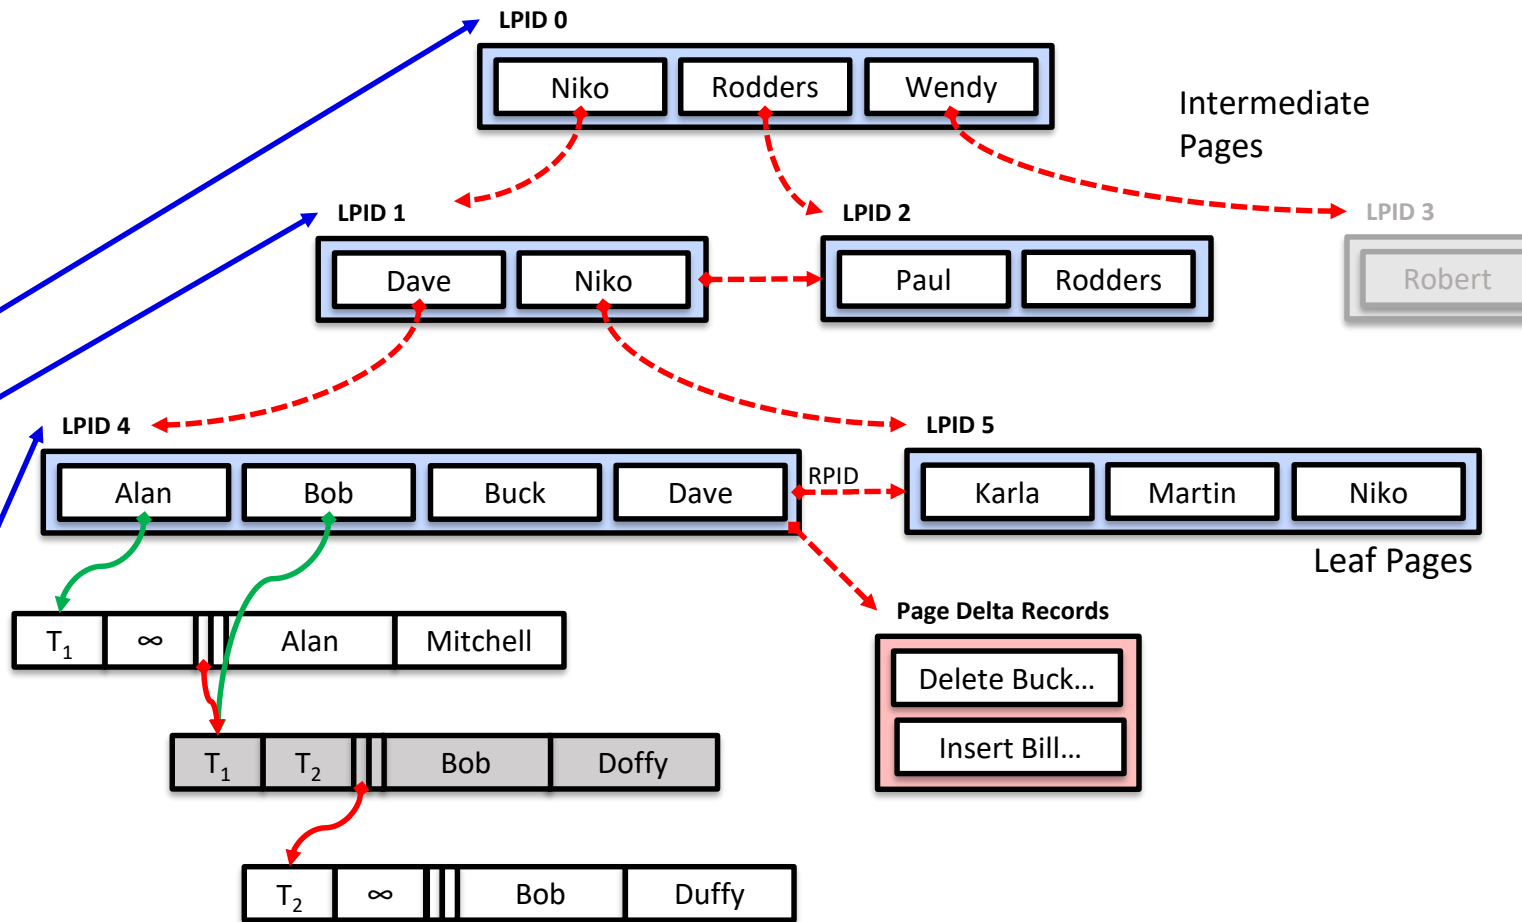
Hash Indexes



BW-Trees

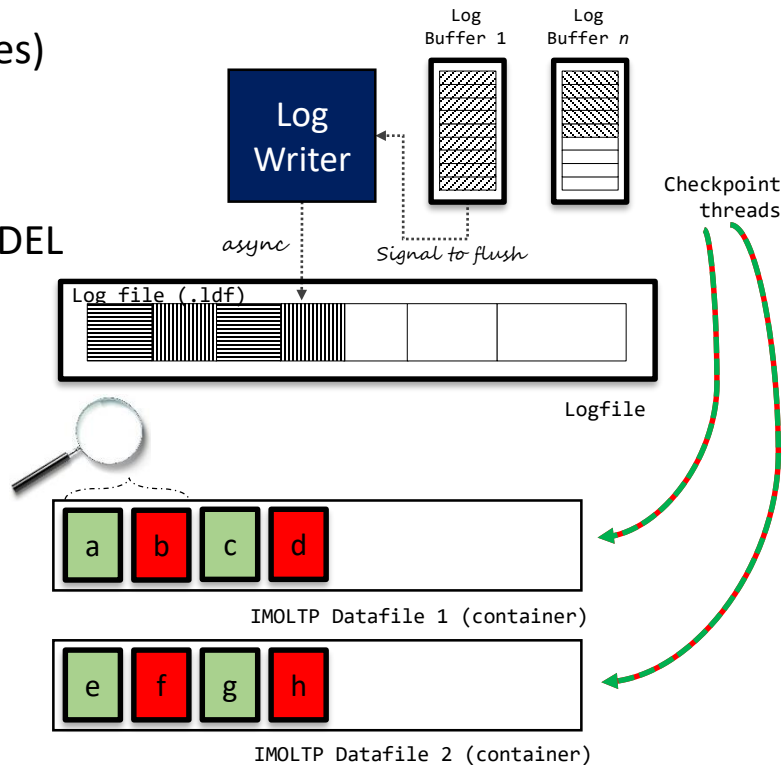

Page mapping table

LPID	PMA
0	0x00008862
1	0x00009862
2	0x00008AA2
3	0x000077C6
4	0x00005B62
5	0x000088BC



Data Persistence

- Checkpoint file pairs, minimum of 8 (in various states)
- Log file primary persistence source
- Database recovery from CFP and Log
- INS to data file, DEL to delta file, Update is INS and DEL



Log Performance improvements

- Indexes not persisted, rebuilt on start-up
 - So NO index maintenance logging
- Log records ordering by Transaction End Timestamps (On disk ordered by LSN)
 - Removes requirement for single log stream to disk per DB (implemented in SQL 2016)
- Transaction consolidation into reduced number of log records
 - Because only committed transactions
 - So NO undo overhead!
- NVDIMM support in Windows 2016!
 - [See Accelerating SQL Server 2016 performance with Persistent Memory in Windows Server 2016](#)

Demo

Logging and improvements

Supported Isolation Levels when accessing IMOLTP tables

Isolation Level	"Bad Dependencies"
<div>*1</div> <div>1</div> <div>READ UNCOMMITTED</div>	Dirty Read, Non-Repeatable Read, Phantoms, Lost Updates
<div>2</div> <div>READ COMMITTED</div>	Non-Repeatable Read, Phantoms , Lost Updates

*1 Politely elevated (mapped to SNAPSHOT) through `MEMORY_OPTIMIZED_ELEVATE_TO_SNAPSHOT = ON`

*2 But is supported for auto-commit transactions

<div>3</div> <div>REPEATABLE READ</div>	Phantoms
<div>4</div> <div>SERIALIZABLE</div>	None
<div>5</div> <div>SNAPSHOT</div>	None (though Causal consistency concerns, lost update prevention and other behaviours)

Querying data

- Interop for:
 - Cross container queries (but introduce synchronisation concerns)
 - Best support
 - But uses legacy engine and incurs unnecessary overhead
- Native Compilation for:
 - Best Performance (but...)
 - Is the most restrictive (and cannot query non in-memory tables)

Cross-Container Transactions

- Cross container transactions therefore are really two internal transactions that are synchronized together.
- Cross container (disk/ in-memory) table transactions are supported but only for:
 - READCOMMITTED + in-memory SNAPSHOT
 - READCOMMITTED + in-memory REPEATABLE/ SERIALIZABLE
 - REPEATABLE/ SERIALIZABLE + in-memory SNAPSHOT
- Are synchronization issues for:
 - SNAPSHOT + ANY In-Memory isolation (so cant do it!)

Demo

Isolation

Gotchas

- Not enough memory to load table causes IMOLTP database stuck in recovery *₁
- Running out of memory no transactions
- Combined size of ACTIVE Checkpoint File Pairs on disk equates to **approximately 2x the size of table that's in Memory** (depending upon frequency of updates)
- In-memory row versions and **maintenance operations can require in excess of 2-3x size of table in Memory**
- Long running transactions consume more memory for in-memory versions
- No data overwritten on disk. Sequential IO is KEY so locate containers on drives optimized for this
- Disk is even more important because:
 - **Transaction log is still king** for transaction speed
 - Log truncation dependant on data persistence!

*₁ Think about database restores

In-Memory OLTP Limitations

Addresses	Feature/Limit	SQL Server 2014	SQL Server 2016
Performance	Parallelism	Not supported	Supported
Performance	Maximum combined size of durable tables	256 GB	2 TB Physical bounds
Performance	Offline Checkpoint Threads	1	1 per container
Performance/ Management	ALTER PROCEDURE / sp_recompile	Not supported	Supported (fully online)
Performance/ Management	ALTER TABLE	Not supported, (DROP / re-CREATE)	Partially supported, (offline operation)
Performance/ Compatibility	Nested native procedure calls	Not supported	Supported
Performance/ Compatibility	Natively-compiled scalar UDFs	Not supported	Supported
Performance/ Compatibility	Indexes on NULLable columns	Not supported	Supported
Compatibility	LOB (varbinary(max), [n]varchar(max))	Not supported	Supported
Compatibility	DML triggers	Not supported	Partially supported (AFTER, natively compiled)
Compatibility	Non-BIN2 collations in index key columns	Not supported	Supported
Compatibility	Non-Latin codepages for [var]char columns	Not supported	Supported
Compatibility	Non-BIN2 comparison / sorting in native modules	Not supported	Supported
Integrity/ Compatibility	Foreign Keys	Not supported	Supported
Integrity/ Compatibility	Check/Unique Constraints	Not supported	Supported
Compatibility	OUTER JOIN, OR, NOT, UNION [ALL], DISTINCT, EXISTS, IN	Not supported	Supported
Compatibility	Multiple Active Result Sets	Not supported	Supported
Management	SSMS Table Designer	Not supported	Supported
Security	Transparent Data Encryption (TDE)	Not supported	Supported




2014 Unsupported features [https://msdn.microsoft.com/en-us/library/dn246937\(v=sql.120\).aspx](https://msdn.microsoft.com/en-us/library/dn246937(v=sql.120).aspx)

2016 Unsupported features [https://msdn.microsoft.com/en-us/library/dn246937\(v=sql.130\).aspx](https://msdn.microsoft.com/en-us/library/dn246937(v=sql.130).aspx)

List based from original post <http://bit.ly/2euFlxz> at SQLPerformance.com written by Aaron Bertrand

Further reading (research papers and books)

- In-Memory OLTP (In-Memory Optimization) <https://msdn.microsoft.com/en-us/library/dn133186.aspx>
- High-Performance Concurrency Control Mechanisms for Main-Memory Databases - Per-Åke Larson, Spyros Blanas, Cristian Diaconu, Craig Freedman, Jignesh M. Patel, Mike Zwilling
- The Hekaton Memory-Optimized OLTP Engine - Per-Ake Larson, Mike Zwilling, Kevin Farlee
- Concurrent Programming Without Locks — Keir Fraser, Tim Harris
- The Bw-Tree: A B-tree for New Hardware Platforms — Justin J. Levandoski, David B. Lomet, Sudipta Sengupta



In Summary (what we have learnt today)...

- Rise in CPU cores, abundance of memory and out of date concurrency model is the reason why in-memory OLTP is the future
- Interleaving and concurrent execution is why we hit severe bottlenecks in pessimistic workloads i.e. Blocking is almost assured
- Snapshot Isolation has no bad dependencies, so is a good fit for our new model (as the new default)
- Adoption will be slow due to some of the complexities (or differences) and restrictions (which are being removed). It is both a development and administrative concern
- Microsoft are 110% committed to this technology

Session Evaluations

Your feedback is important and valuable.

3

ways to access

Submit by 5pm
Friday November 6th to
WIN prizes



Go to passSummit.com



Download the GuideBook App
and search: **PASS Summit
2016**



Follow the QR code link displayed
on session signage throughout the
conference venue and in the
program guide

Thank you for listening!

Email: mark.broadbent@sqlcambs.org.uk

Twitter: [retracement](#)

Blog: <http://tenbulls.co.uk>

Slideshare: <http://www.slideshare.net/retracement>